

# Summary of Research and Proposals

DSP Lab 2017

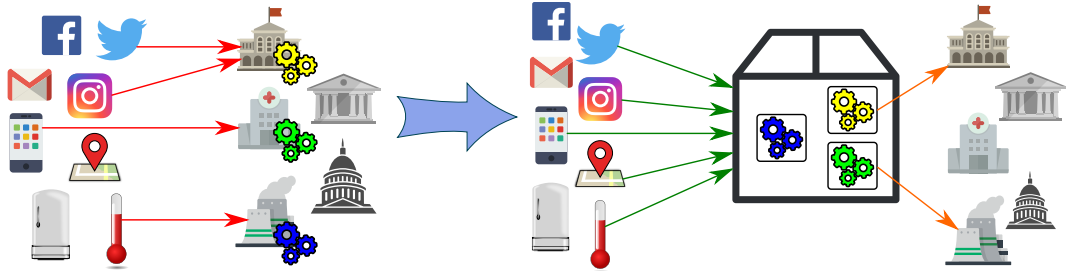
Yousef Amar



2017-10-02

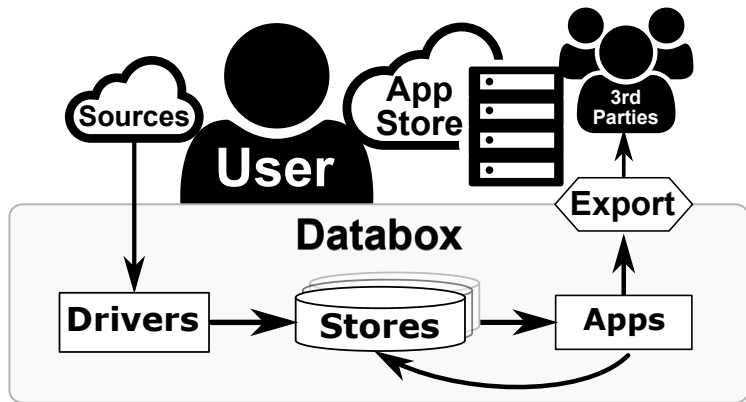
# Research Context

## The Databox Platform



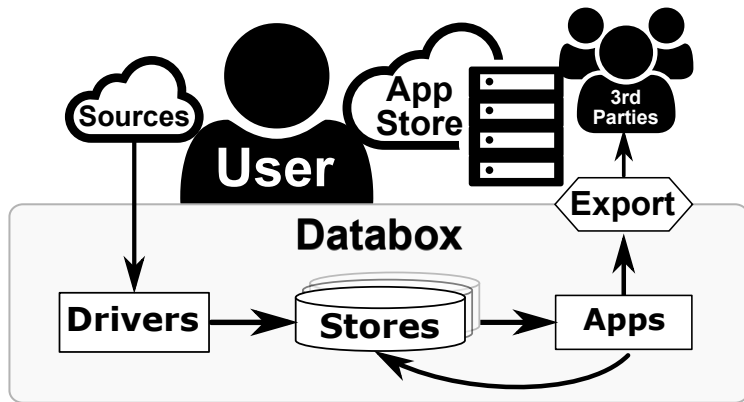
# Research Context

## The Databox Platform



## Research Context

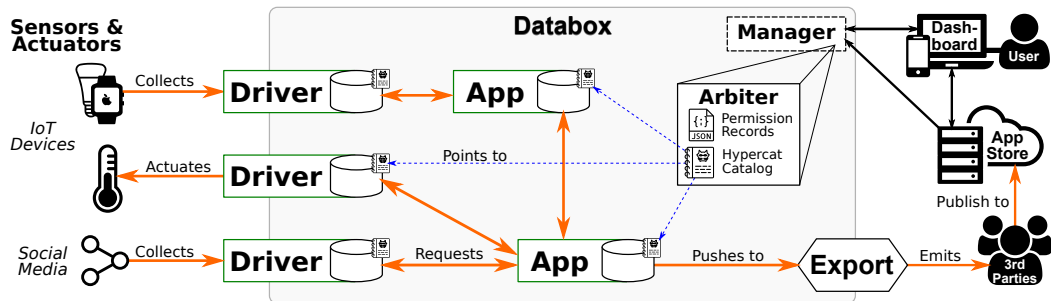
The Databox Platform



*How can we design safe, scalable access control systems with arbitrary restrictions in this context?*

# Research Context

## The Databox Platform



# Implementation

## The Route

- ▶ Triad of *target*, *path*, and *method*
- ▶ The container as a host
- ▶ RESTful APIs for all operations
- ▶ Direct mapping of HTTP methods to CRUD functions
- ▶ Per-route granular permissions

```
{  
  "target": "smartphone-store",  
  "path": "/accelerometer/ts/latest",  
  "method": "POST"  
}  
  
{  
  "target": "smartphone-store",  
  "path": "/(sub|unsub)/gps/*",  
  "method": "GET"  
}
```

# Implementation

## Delegated Authorization

- ▶ Google Research: Macaroons
  - ▶ A standard similar to signed cookies
  - ▶ Can be attenuated by “caveats”
  - ▶ Embedded permissions
  - ▶ Minting and verification can be separated through shared secret keys

```
target = smartphone-store  
path = /(sub|unsub)/gps/*  
method = GET  
time < 1489405851417
```

```
target = smartphone-store  
path = /light/ts/range  
method = GET  
startTimestamp >= 1489405234352  
endTimestamp <= 1489405259525
```



# Implementation

## Resource Discovery

- ▶ API for describing APIs
- ▶ Directory servers
- ▶ Many competing standards
  - ▶ Resource Description Framework (RDF)
  - ▶ Web Application Description Language (WADL)
  - ▶ Web Services Description Language (WSDL)
  - ▶ eXtensible Resource Descriptor (XRD)
- ▶ Subject-predicate-object style prevalent
- ▶ Different formats and applications — XML for REST, SOAP, OpenID



# Implementation


## Resource Discovery

- ▶ Hypercat: Recently joined BSI Group
- ▶ IoT-first specification design
- ▶ JSON/REST over XML/SOAP
- ▶ Only cataloguing; ontologies and authorisation extensible
- ▶ Discoverability vs accessibility
- ▶ Catalogues can be nested, allowing decentralisation and distribution

```
{
  "catalogue-metadata": [{
    "rel": "urn:X-hypercat:rels:isContentType",
    "val": "application/vnd.hypercat.catalogue+json"
  }, {
    "rel": "urn:X-hypercat:rels:hasDescription:en",
    "val": "A Databox Store"
  }],
  "items": [{
    "href": "http://some-store/light",
    "item-metadata": [{
      "rel": "urn:X-hypercat:rels:hasDescription:en",
      "val": "Light Datasource"
    }, {
      "rel": "urn:X-databox:rels:hasVendor",
      "val": "Databox Inc."
    }, {
      "rel": "urn:X-databox:rels:isActuator",
      "val": false
    }
  ]
}]
}
```

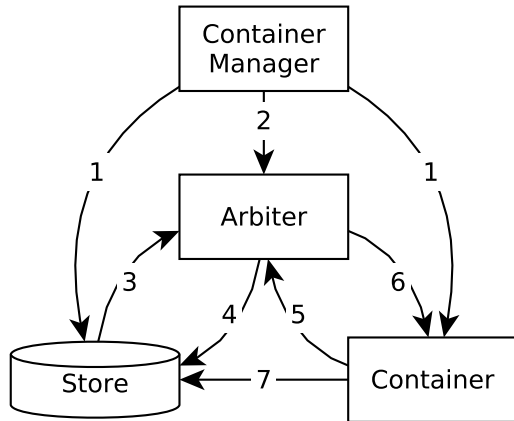
# Implementation

## The Arbiter



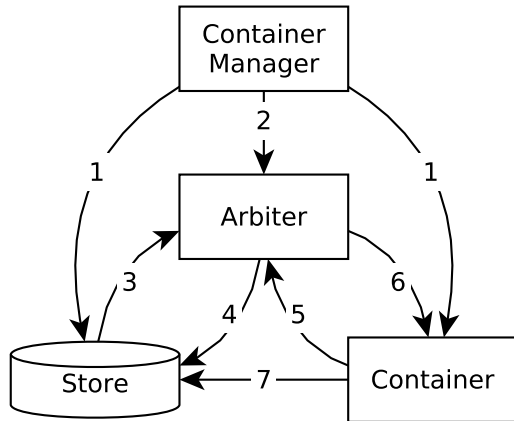
# Implementation

## Authorisation Flow



# Implementation


## Authorisation Flow



1. CM passes unique tokens

# Implementation

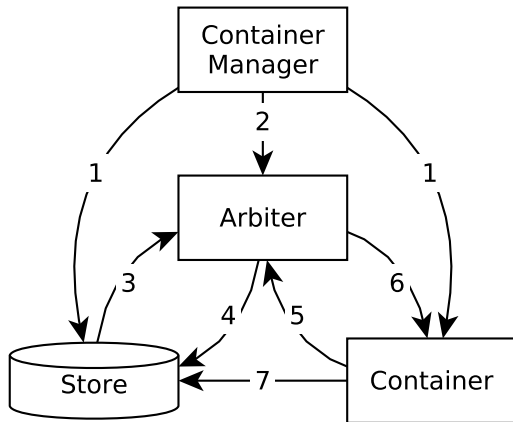
## Authorisation Flow



1. CM passes unique tokens
2. CM updates permissions

# Implementation

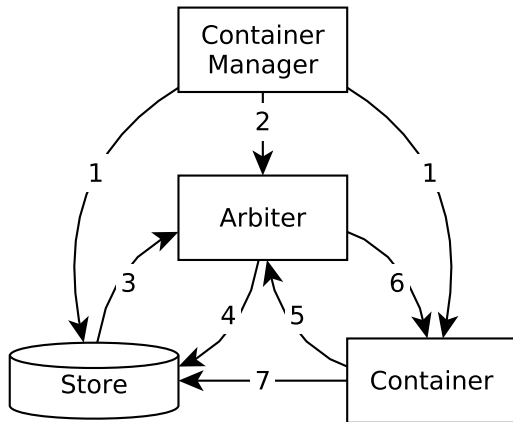
## Authorisation Flow



1. CM passes unique tokens
2. CM updates permissions
3. Store registers itself

# Implementation

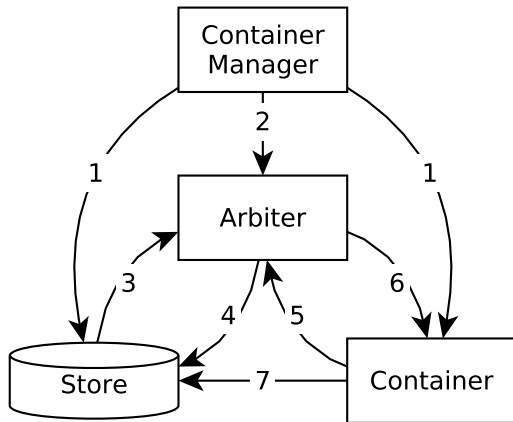
## Authorisation Flow



1. CM passes unique tokens
2. CM updates permissions
3. Store registers itself
4. Arbiter responds with shared secret

# Implementation

## Authorisation Flow

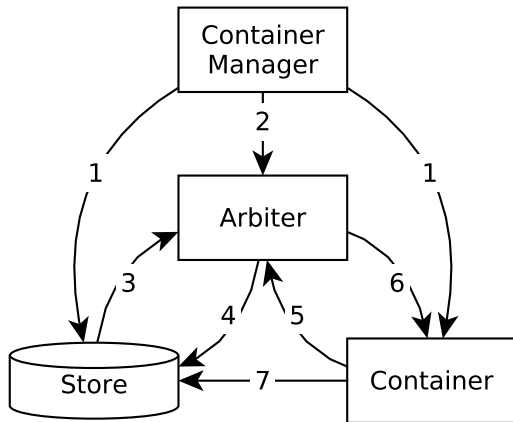


1. CM passes unique tokens
2. CM updates permissions
3. Store registers itself
4. Arbiter responds with shared secret
5. Container requests bearer token



# Implementation

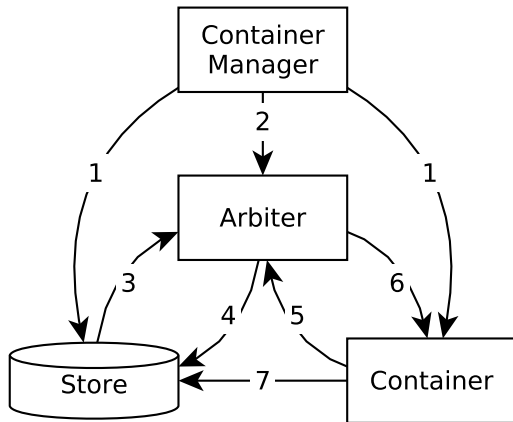
## Authorisation Flow



1. CM passes unique tokens
2. CM updates permissions
3. Store registers itself
4. Arbiter responds with shared secret
5. Container requests bearer token
6. Arbiter checks and responds

# Implementation

## Authorisation Flow




1. CM passes unique tokens
2. CM updates permissions
3. Store registers itself
4. Arbiter responds with shared secret
5. Container requests bearer token
6. Arbiter checks and responds
7. Container can now read/write to store

# Implementation

## Transcription of Permissions

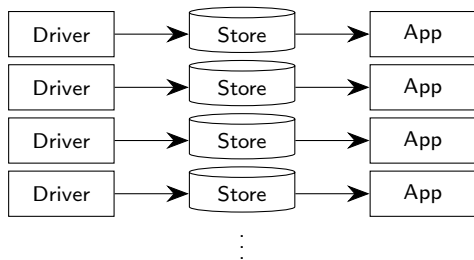
1. Drivers/apps come packaged with a *manifest*
  - ▶ Contain image metadata
  - ▶ Enumerate granular permissions for sources, concurrency, external access, and hardware
2. Users generate a Service-level Agreement (SLA)
3. The arbiter records granted permissions
4. Tokens are minted based on these



```
{  
  "name": "app",  
  "author": "amar",  
  "permissions": [  
    {  
      "source": "twitter",  
      "required": true  
    },  
    {  
      "source": "gps"  
    },  
    {},  
    {}  
  ]  
}
```

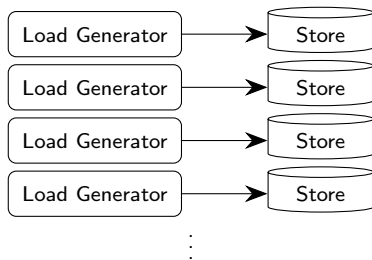
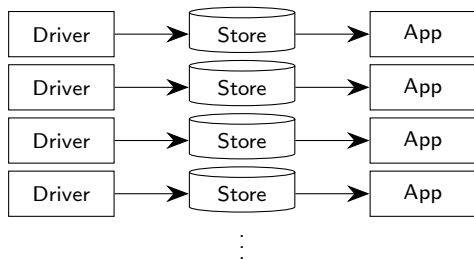
# Evaluation

## Scalability



# Evaluation

## Scalability



# Evaluation

## Scalability

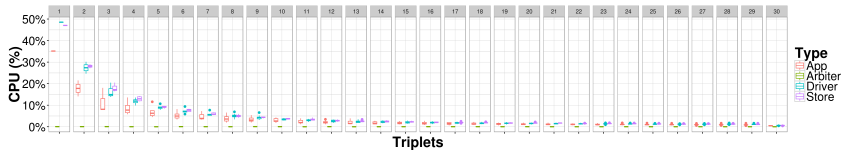


Figure: Percentage CPU Usage by Container Type

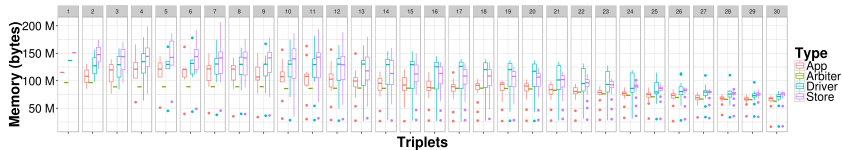


Figure: Memory Usage by Container Type

# Evaluation

## Scalability

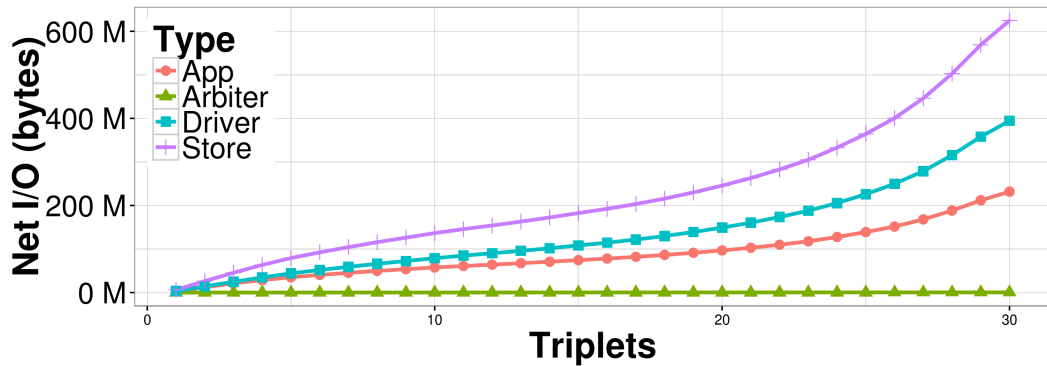


Figure: Sum Net I/O by Container Type

# Evaluation

## Scalability

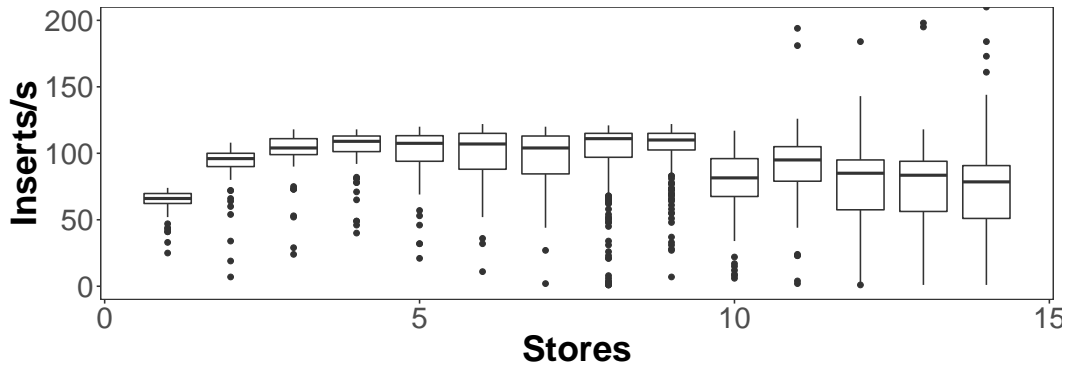


Figure: Inserts/s over Stores under Maximum Load



# Evaluation

## Scalability

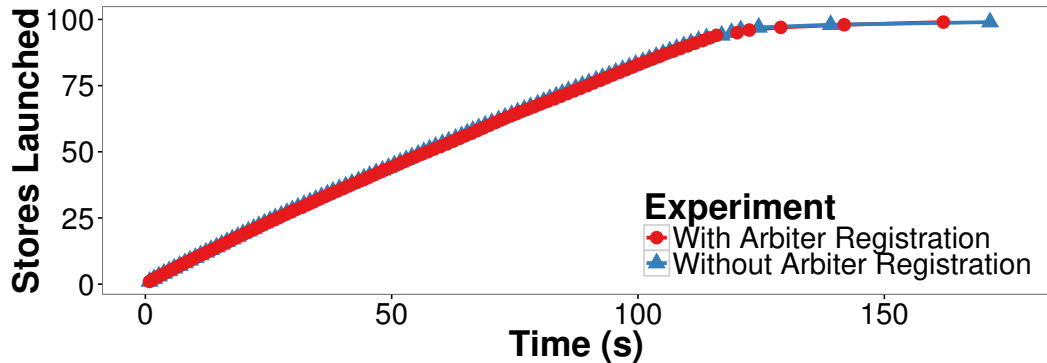


Figure: Stores Launched over Time

# Evaluation

## Topology

Differences in Time to Availability (TTA)

1. Device  $\rightarrow$  Cloud:  
 $65ms$
2. Device  $\rightarrow$  Cloud  $\rightarrow$  Home:  
 $83ms$
3. Device  $\rightarrow$  Home:  
 $78ms$
4. Device  $\rightarrow$  Home  $\rightarrow$  Cloud:  
 $80ms$

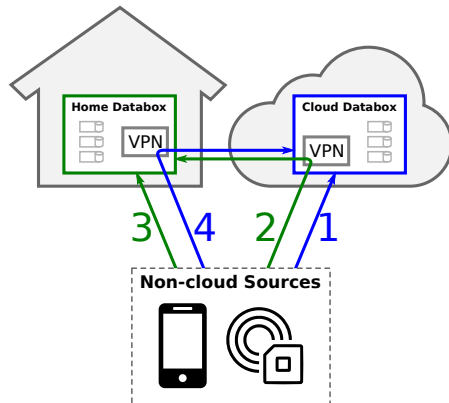


Figure: The four possible data flow scenarios tested

# Evaluation

## Topology

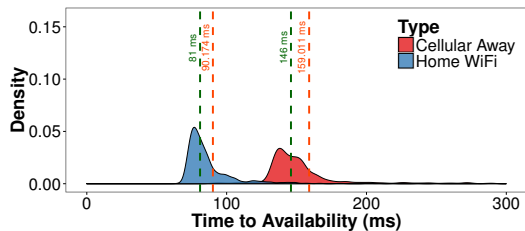


Figure: Data Time to Availability from Device to Cloud Databox Directly




Figure: Data Time to Availability from Device to Home Databox Directly

# Evaluation

## Topology

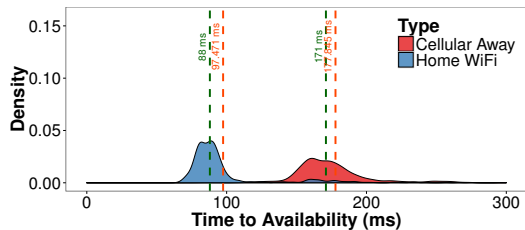


Figure: Data Time to Availability from Device to Home Databox via Cloud VPN




Figure: Data Time to Availability from Device to Cloud Databox via Home VPN

# Evaluation

## Topology

- ▶ TTA source away from home  $>$  source at home
- ▶ So minor, barely indistinguishable from NTP drift
- ▶ Based on performance alone, UX indifferent
- ▶ Scenarios through home (especially when source is away) have mean shifted right due to latency spikes
- ▶ Direct connections mean lower TTA, and cloud faster than home *ceteris paribus*
- ▶ Small difference for devices as sources vs cloud servers
- ▶ For devices, processing at home  $>$  in the cloud  $\pm$  NTP error even ignoring privacy advantages
- ▶ Home vs cloud — reliability vs cost
- ▶ Pure cloud only more advantageous for off-site processing (e.g. GPU-heavy image processing)

# Evaluation

Time to Availability

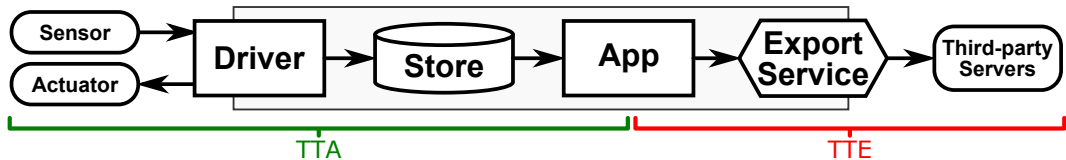


Figure: Sections of the data pipeline timed

# Evaluation

## Time to Availability

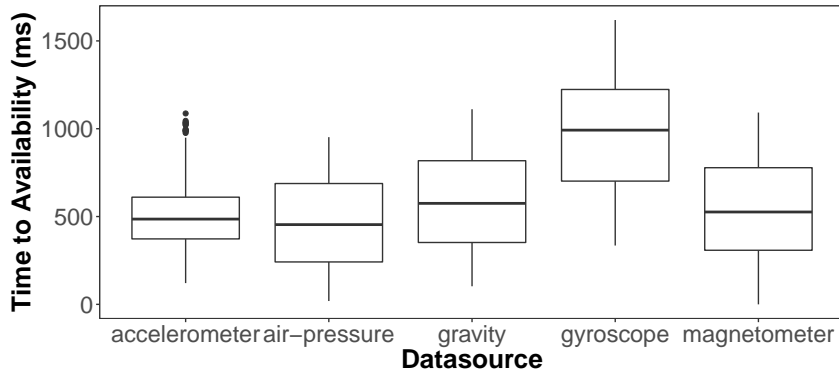


Figure: Time-to-Availability (TTA) on a Raspberry Pi for high-frequency sensors

# Evaluation

## Time to Availability

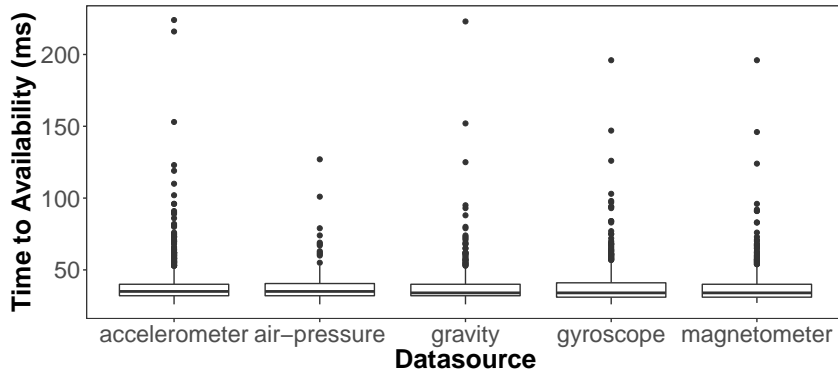
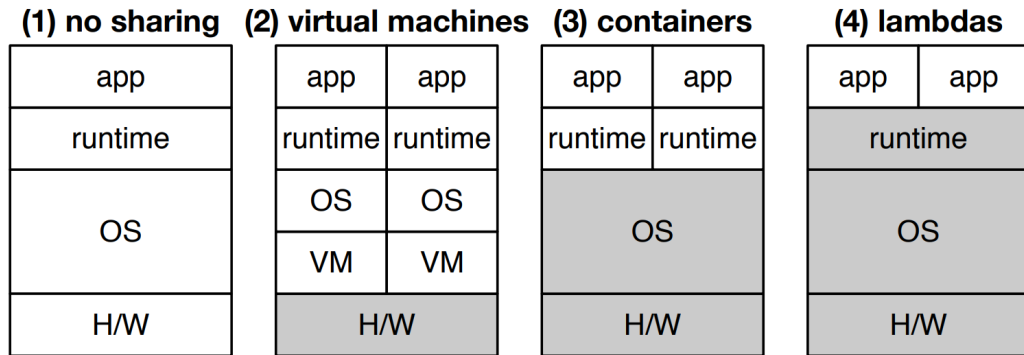


Figure: Time-to-Availability (TTA) on an Intel NUC for high-frequency sensors



# The Serverless Paradigm

## Background



**Figure 1: Evolution of Sharing.** *Gray layers are shared.*

**Figure:** Hendrickson, et al. "Serverless computation with openlambda." Elastic 60 (2016): 80.

# The Serverless Paradigm Architecture

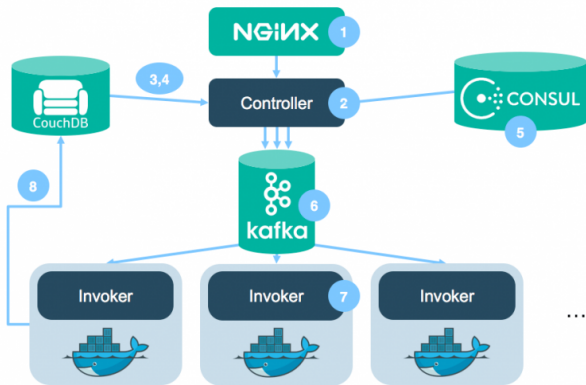


Figure: IBM's High-level OpenWhisk Architecture Diagram

# The Serverless Paradigm

## Architecture

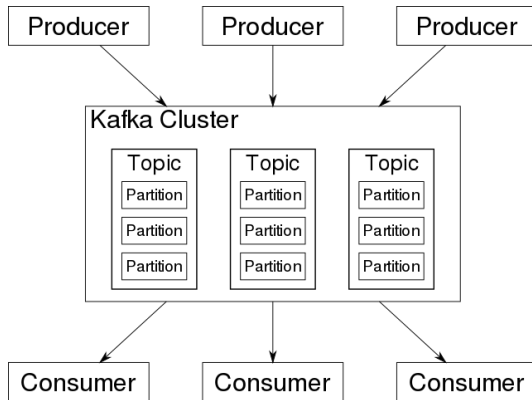
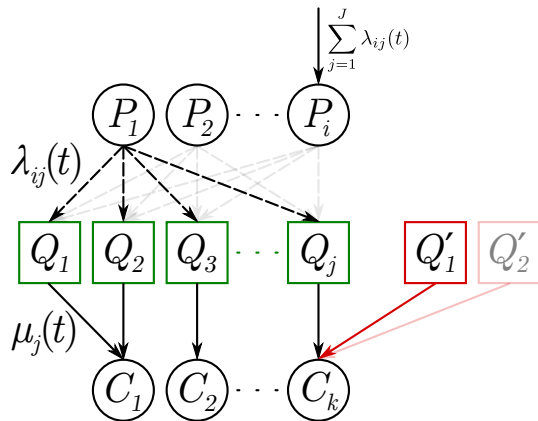


Figure: Apache Kafka High-level Architecture Diagram

# Low-latency Serverless Approach



**Figure:** An Overview of Inter-component Relationships

## Low-latency Serverless Approach

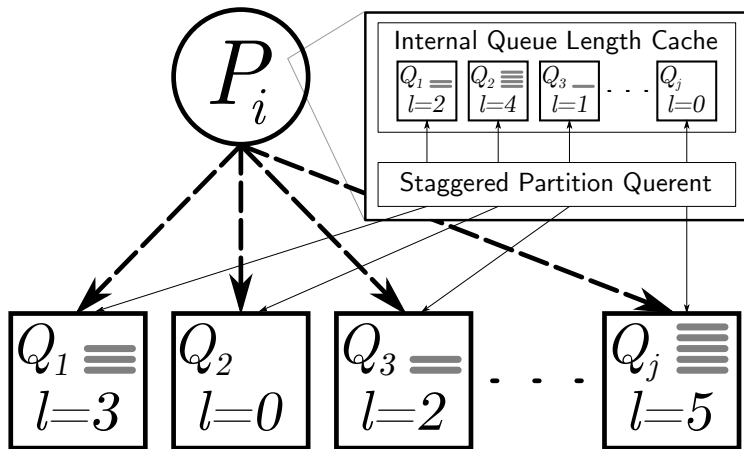



Figure: The Internal Components of a Producer

# Low-latency Serverless Approach



**Figure:** Producer-intrinsic Steps for Computing Partitioning Decisions from Stale Queue Lengths

## Low-latency Serverless Simulations

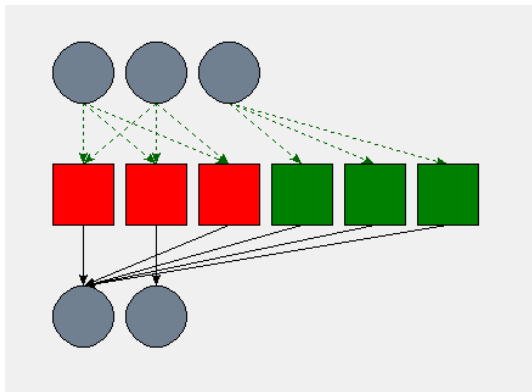


Figure: An Example of Simulation Topology

# Low-latency Serverless Simulations





Figure: Simulation Results with Different Partitioning Algorithms



# Next Steps

## Privacy and Risk Metrics





Figure: One Proof of Concept Experiment – Surprisal over Real Smart Meter Data

# Next Steps

## Serverless over Transient Clouds

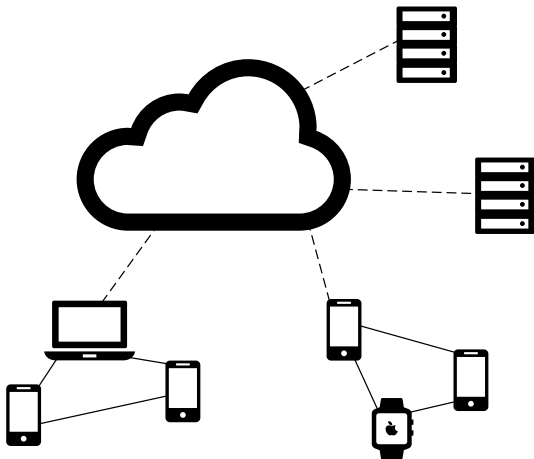
- ▶ Serverless on the edge
- ▶ Optimising for context through latency
- ▶ Processor selection based on arbitrary metrics, e.g. surprisal



# Next Steps

## Transient Privacy-Aware Clouds

- ▶ Encoding user-defined thresholds into bearer tokens
- ▶ Joint context at hierarchical levels
- ▶ TCACs → TPACs?



# Next Steps

## The Big Picture

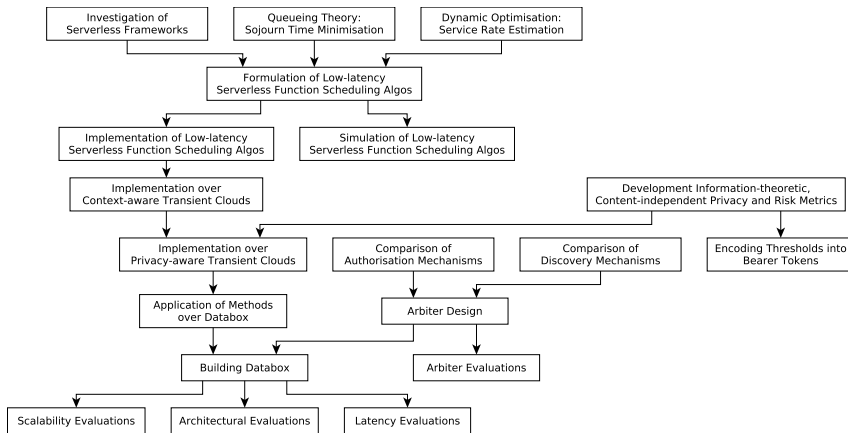


Figure: A High-level Dependency Graph of Research Activities

# Thank you for your attention!

Questions?

More info: <http://yousefamar.com/>